

# LectureNote2: CS221 (FALL2021, Stanford)

## Linear Regression

### *The discovery of Ceres*

- 1801: astronomer Piazzi discovered Ceres, made 19 observations of location before it was obscured by the sun
- September 1801: Gauss took Piazzi's data and created a model of Ceres's orbit, makes prediction
- December 7, 1801: Ceres located within 1/2 degree of Gauss's prediction, much more accurate than other astronomers
  - METHOD: Least squares linear regression

### *Linear Regression Framework*

$D_{train} - learningalgorithm \rightarrow prediction[f]$

$x \quad y$

1    1

2    3

4    3

- Design decisions:
  - i. Which predictors are possible? *hypothesis class*
  - ii. How good is a predictor? *loss function*
  - iii. How do we compute the the best predictor? *optimization algorithm*

- Hypothesis Class: which predictors?

- $f(x) = w_1 + w_2x$

- Vector Notation

- $\vec{w}_{weight\ vector} = [w_1, w_2]$

- $\phi_{feature\ extractor} = [1, x]_{feature\ vector}$

- $f_{\vec{w}}(x) = \vec{w} \cdot \phi(x)$

- Hypothesis Class:  $F = \{ f_{\vec{w}}(x) : \vec{w} \in \mathbb{R}^2 \}$
- Loss Function: how good is a predictor?
  - $Loss(x, y, \vec{w}) = (f_{\vec{w}}(x) - y)^2$  *squared loss*
    - residual:  $|f_{\vec{w}}(x) - y|$
  - $TrainLoss(\vec{w}) = \frac{1}{|D_{train}|} \sum_{(x,y) \in D_{train}} Loss(x, y, \vec{w})$
- Optimization Algorithm: how to compute best?
  - Goal:  $\min_{\vec{w}} TrainLoss(\vec{w})$
  - Definition: gradient
    - The gradient  $\nabla_{\vec{w}} TrainLoss(\vec{w})$  is the direction that increases the training loss the most
  - Algorithm: gradient descent
    - Initialize  $\vec{w} = [0, \dots, 0]$
    - For  $t = 1, \dots, T$ : epochs
      - $\vec{w} \leftarrow \vec{w} - \eta_{step\ size} \nabla_{\vec{w}} TrainLoss(\vec{w})$
    - initial weight = 0
    - we want to find the weight vector that minimizes the value of TrainLoss function
    - update weight vector by going to the direction that decreases the TrainLoss function's value by subtracting gradient
- Computing the gradient
  - $\nabla_{\vec{w}} TrainLoss(\vec{w}) = \frac{1}{|D_{train}|} \sum_{(x,y) \in D_{train}} 2(\vec{w} \cdot \phi(x) - y) \phi(x) = \frac{1}{|D_{train}|} \sum_{(x,y) \in D_{train}} 2(residual) \phi(x)$
  - if (prediction = target)  $\rightarrow$  gradient = 0
  - if gradient at the end of iteration be 0, it means the gradient descent has converged


# Linear Classification

training data

$x_1$	$x_2$	$y$
0	2	1
-2	0	1
1	-1	-1


Learning Algorithm

$[2, 0]$  input



classifier

-1 label



- Design decisions:
    - i. Which classifiers are possible? *hypothesis class*
    - ii. How good is a classifier? *loss function*
    - iii. How do we compute the best classifier? *optimization algorithm*
  - An example linear classifier
    - $f(\vec{x}) = \text{sign}([-0.6, 0.6] \cdot [x_1, x_2]) = \text{sign}(\vec{w} \cdot \phi(\vec{x}))$
    - $\text{sign}(z) = \begin{cases} +1 & (z > 0) \\ -1 & (z < 0) \\ 0 & (z = 0) \end{cases}$
    - if the angle between  $\vec{w}$  and  $\phi(\vec{x})$  is **acute**, the sign of the dot product of these two is +
    - if the angle between  $\vec{w}$  and  $\phi(\vec{x})$  is **obtuse**, the sign of the dot product of these two is -
    - if the angle between  $\vec{w}$  and  $\phi(\vec{x})$  is **perpendicular**, the dot product of these two is 0  $\rightarrow$
- Decision Boundary**
- $\therefore$  linear classifier labels the inputs according to the angle of  $\phi(\vec{x})_{\text{input}}$  with respect to  $\vec{w}_{\text{weight vector}}$

- Hypothesis Class: which classifiers?

- General binary classifier:  $f_{\vec{w}}(\vec{x}) = \text{sign}(\vec{w} \cdot \phi(x))$
- Hypothesis class:  $F = \{f_{\vec{w}}(\vec{x}) : \vec{w} \in \mathbb{R}^2\}$

- Loss Function: how good is a classifier?


- $Loss_{0-1}(\vec{x}, y, \vec{w}) = 1[f_{\vec{w}}(\vec{x}) \neq y]$  *zero-one loss*
  - $Loss_{0-1}(\vec{x}, y, \vec{w}) = \begin{cases} 1 & (f_{\vec{w}}(\vec{x}) \neq y) \\ 0 & (f_{\vec{w}}(\vec{x}) = y) \end{cases}$

- Score and Margin

- predicted label:  $f_{\vec{w}}(\vec{x}) = \text{sign}(\vec{w} \cdot \phi(\vec{x}))$
- target label:  $y$
- Definition: score
  - The score on an example  $(x, y)$  is  $\vec{w} \cdot \phi(\vec{x})$ , how **confident** we are in predicting
  - *raw output of the model before applying a decision threshold*
- Definition: margin
  - The margin on an example  $(x, y)$  is  $(\vec{w} \cdot \phi(\vec{x}))y$ , how **correct** our prediction is
  - *positive*  $\rightarrow$  *correct*
  - *negative*  $\rightarrow$  *incorrect*
  - $|\text{margin}| = \text{distance from the decision boundary}$
- $Loss_{0-1}(\vec{x}, y, \vec{w}) = 1[f_{\vec{w}}(\vec{x}) \neq y] = 1[\text{margin} \leq 0] = \begin{cases} 1 & (\text{margin} \leq 0) \\ 0 & (\text{margin} > 0) \end{cases}$

- Optimization Algorithm: how to compute best?


- $\nabla_{\vec{w}} \text{TrainLoss}(\vec{w}) = \sum_{(x,y) \in D_{\text{train}}} \nabla Loss_{0-1}(\vec{x}, y, \vec{w})$
- $\nabla Loss_{0-1}(\vec{x}, y, \vec{w}) = \nabla[(\vec{w} \cdot \phi(\vec{x})) y \leq 0]$
- $\nabla Loss_{0-1}(\vec{x}, y, \vec{w})$  is zero almost everywhere!  $\rightarrow$  we can't use gradient descent




- Solutions

- a. Hinge loss

- $Loss_{hinge}(\vec{x}, y, \vec{w}) = \max\{1 - margin, 0\}$



- The 1 can be any positive number, and the magnitude determines the regularization strength
  - Hinge loss is the upper bound on the zero-one loss
  - $\nabla_{\vec{w}} Loss_{hinge}(\vec{x}, y, \vec{w}) = \begin{cases} -\phi(\vec{x}) y & (1 > margin) \\ 0 & (otherwise) \end{cases}$
  - Even if an input is classified correctly, it can still generate loss if it didn't meet the threshold (in this case, 1)
- b. Logistic Regression
- $Loss_{hinge}(\vec{x}, y, \vec{w}) = \log(1 + e^{-margin})$



- Intuition: try to increase margin even when it already exceeds 1

# Summary

$\Phi$ score: $\vec{w} \circ \phi(x)$	Regression	Classification
Prediction $f_{\vec{w}}(x)$	score	sign of score
Relate to target $y$	residual (score - $y$ )	margin (score * $y$ )
Loss functions	Squared Absolute deviation	zero-one hinge logistic
Algorithm	gradient descent	