

LectureNote2: CS229 (SPRING2022, Stanford)


Supervised Learning

Set Up

- Prediction:

$$\begin{array}{ccc} x & \rightarrow & y \\ \circ \ h : & \begin{array}{c} \textit{images} \\ \textit{text} \\ \textit{housedata} \end{array} & \begin{array}{c} \textit{cat} \\ \textit{hate speech?} \\ \textit{price} \end{array} \end{array}$$

- Given: training set
 - $\{ (x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}) \}$
 - y = supervision
- Do: find **good** $h(\textit{hypothesis}) : x \rightarrow y$
 - we care about **new** x 's that are not in our training set
 - if y is *discrete* \rightarrow Classification
 - if y is *continuous* \rightarrow Regression
- How do we represent h ?
 - $h(x) = \theta_0 + \theta_1 x_1$



- Generalization

	$size(x_1)$	$bedroom(x_2)$	$lot\ size(x_3)$...	$price(y)$
$x^{(1)}$	2104	4	45k	...	400
$x^{(2)}$	2500	3	30k	...	900

- $$h(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_d x_d = \sum_{j=0}^d \theta_j x_j \quad (x_0 = 1)$$

- $$\theta_{parameters} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix} \quad x_{features}^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_d^{(i)} \end{bmatrix}$$

- $(x^{(i)}, y^{(i)}) \leftarrow$ i-th training example

- $$X_{train\ data} = \begin{bmatrix} -x^{(1)}- \\ -x^{(2)}- \\ \vdots \\ -x^{(n)}- \end{bmatrix} \in \mathbb{R}^{n(d+1)}$$

- +1 is the extra dimension x_0 , whose value is set to be 1 above

- $h_{\theta}(x) = \sum_{j=0}^d \theta_j x_j$ ($x_0 = 1$), I want $h_{\theta}(x) \simeq y$
- $J(\theta)_{cost\ function} = \frac{1}{2} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2_{least\ squares}$
 - $\frac{1}{2}$ is for convenience purpose (It cancels out with 2 when we take the derivative)
 - we only care about the minimizer, which is the θ that minimizes $J(\theta)$. This process is called "Optimization"
- Gradient Descent
 - $\theta^{(0)} = 0$
 - $\theta_j^{(t+1)} = \theta_j^{(t)} - \alpha \frac{\partial}{\partial \theta_j} J(\theta^{(t)})$ (α = learning rate, $j = 0 \dots d$)
 - $\frac{\partial}{\partial \theta_j} J(\theta^{(t)}) = \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) \frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)})$
 - $h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_d x_d$, $\frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)}) = x_j^{(i)}$
 - $\therefore \theta_j^{(t+1)} = \theta_j^{(t)} - \alpha \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$
 - i-th data point, j-th component
 - for all j 's, $\theta^{(t+1)} = \theta^{(t)} - \alpha \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$
 - α can be changed throughout the iteration
- Batch vs Stochastic Minibatch
 - Minibatch: randomly select a B ($B \ll n$)
 - calculate a noisy estimation $\theta^{(t+1)} = \theta^{(t)} - \alpha_B \sum_{i \in B} (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$